

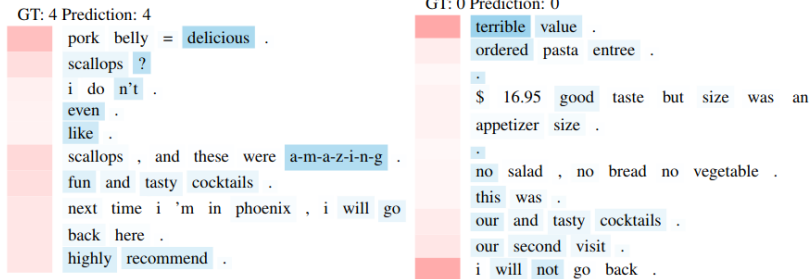
# Attention Mechanism for Neural Networks

Toby Dylan Hocking

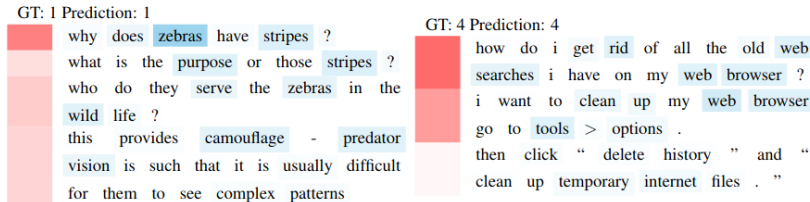
# Supervised machine learning

- ▶ Goal is to learn a function  $f(\mathbf{x}) = y$  where  $\mathbf{x} \in \mathbb{R}^p$  is an input/feature vector and  $y$  is an output/label.
- ▶  $\mathbf{x}$  = text document (email, movie review), binary classification  $y \in \{1, 0\}$  (spam or not, good or bad).
- ▶ Last week we studied recurrent neural networks (RNNs), which can be used to model the sequential nature of text data.
- ▶ This week we will augment the RNNs with an attention mechanism, that allows the neural network to prioritize which items in the sequence are important for prediction.

# Intuition/visualization of attention weights



**Figure 5:** Documents from Yelp 2013. Label 4 means star 5, label 0 means star 1.



**Figure 6:** Documents from Yahoo Answers. Label 1 denotes Science and Mathematics and label 4 denotes Computers and Internet.

Yang et al, 2016. Hierarchical Attention Networks for Document Classification.

## A modification of the recurrent neural net from last week

- ▶ Last week we used the final hidden state as the features for predicting the output.
- ▶ In math notation assume we have a sequence of  $\tau$  inputs  $x_1, \dots, x_\tau$ .
- ▶ The RNN  $r$  gives us  $r(x_t) = o_t, h_t$ , output and hidden state.
- ▶ We learned a function  $f(h_\tau)$  for predicting the output label  $y$  of that sequence using the last hidden state.
- ▶ Instead we could use an average of the hidden states  $h_t$ , say  $\sum_{t=1}^{\tau} h_t / \tau$ .
- ▶ What if we could learn the weights in this average, instead of giving each item in the sequence a uniform weight of  $1/\tau$ ?
- ▶ That is the main idea of the attention mechanism.

## Basic idea of attention mechanism

- ▶ We learn a nonlinear transform of the hidden state features,  $u_t = \phi(h_t)$  (single Linear layer followed by nonlinear activation).
- ▶ We learn a weight vector  $w$  which is used to compute normalized importance weights  $\alpha_t = \exp(w^T u_t) / \sum_j \exp(w^T u_j)$  (softmax results in weights that sum to 1).
- ▶ We use the normalized importance weights to average the hidden states,  $\sum_{t=1}^T h_t \alpha_t$ .
- ▶ We can visualize the  $\alpha_t$  values to see which words are important.